

Eval val-PPL vs tokens: Faithful vs IMU-1 (both COMPLETE @ 1.19B tok)
src: qwen3_baseline2tpp_train.log (eval lines + after.txt) ; qwen3_imu1_2tpp_train.log (eval lines)
step->tokens via tok/step=65,536 (verified in both logs); init PPL ~185k omitted (off-scale)

